

CS 7641 - Machine Learning Assignment 1 -- Minming Zhao (gtID: 902685774)

1. Problem Statement

In this report, we will solve two binary classification problems for two datasets using R.

Problem #1: Student alcohol addiction prediction

To study student possibility to be alcohol addictive with relation to their school, gender, age, address, parents' job, their study time etc. Knowing better about student alcohol addiction could help families and teachers pay more attentions to those students who are more likely to be alcohol addictive and help students to study hard and let them enjoy other beautiful aspects of the life other than alcohol addiction.

Problem #2: Adult census income prediction

Money making is always an interesting question, isn't it? This is to predict whether one could make income more than \$50k/y or not based on census data. With this study prediction, it help us understand how to make more money by locating some key features who did make over \$50k/y.

2. Table of Content

1. Student alcohol addiction prediction

- a. Dataset explanation
- b. Learning curve for train data size
- c. Modeling using
 - i. Decision tree
 - ii. Neural Networks
 - iii. Boosting on decision tree
 - iv. SVM
 - v. kNN

2. Adult census income dataset

- a. Dataset explanation
- b. Learning curve to for train data size
- c. Modeling using
 - i. Decision tree
 - ii. Neural Networks
 - iii. Boosting on decision tree
 - iv. SVM
 - v. kNN

3. Conclusion

3. Problem #1 Student alcohol addiction prediction

3.1. Dataset Attributes

Features	Attributes
school	binary: school 'GP' or School 'MS'
sex	binary: 'F' - female or 'M' - male
age	numeric: from 15 to 22
address	binary: 'U' - urban or 'R' - rural
famsize	binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3
Pstatus	binary: 'T' - parents living together or 'A' - apart
Medu	numeric: 0 - none, 1:- primary education (4th grade), 2:5th to 9th grade, 3 secondary education or 4: higher education
Fedu	numeric: 0 - none, 1 - primary education (4th grade), 2 : 5th to 9th grade, 3 : secondary education or 4 : higher education
Mjob	nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other'
Fjob	nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other'
traveltime	home to school. numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour
studytime	numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours
failures	numeric: n if $1 \leq n < 3$, else 4
higher	higher education wants. binary: yes or no
internet	binary: yes or no
romantic	a romantic relationship. binary: yes or no

famrel	family relation. numeric: from 1 - very bad to 5 - excellent
freetime	numeric: from 1 - very low to 5 - very high
goout	numeric: from 1 - very low to 5 - very high
health	numeric: from 1 - very bad to 5 - very good
absences	numeric: from 0 to 93
Grades	numeric: from 0 to 20, output target
Consumption_addictive	Alcohol addiction. binary: yes or no

3.2. Learning curve to help decide train data size

We take a look at the learning curve to see what the minimum training size for each algorithm is. As we could notice from the below that SVM require least amount of training size while neural network, boosting and knn will require most training set. For sanity to obtaining best model, we later randomly assign 80% of all data as training set and remaining 20% for test set.

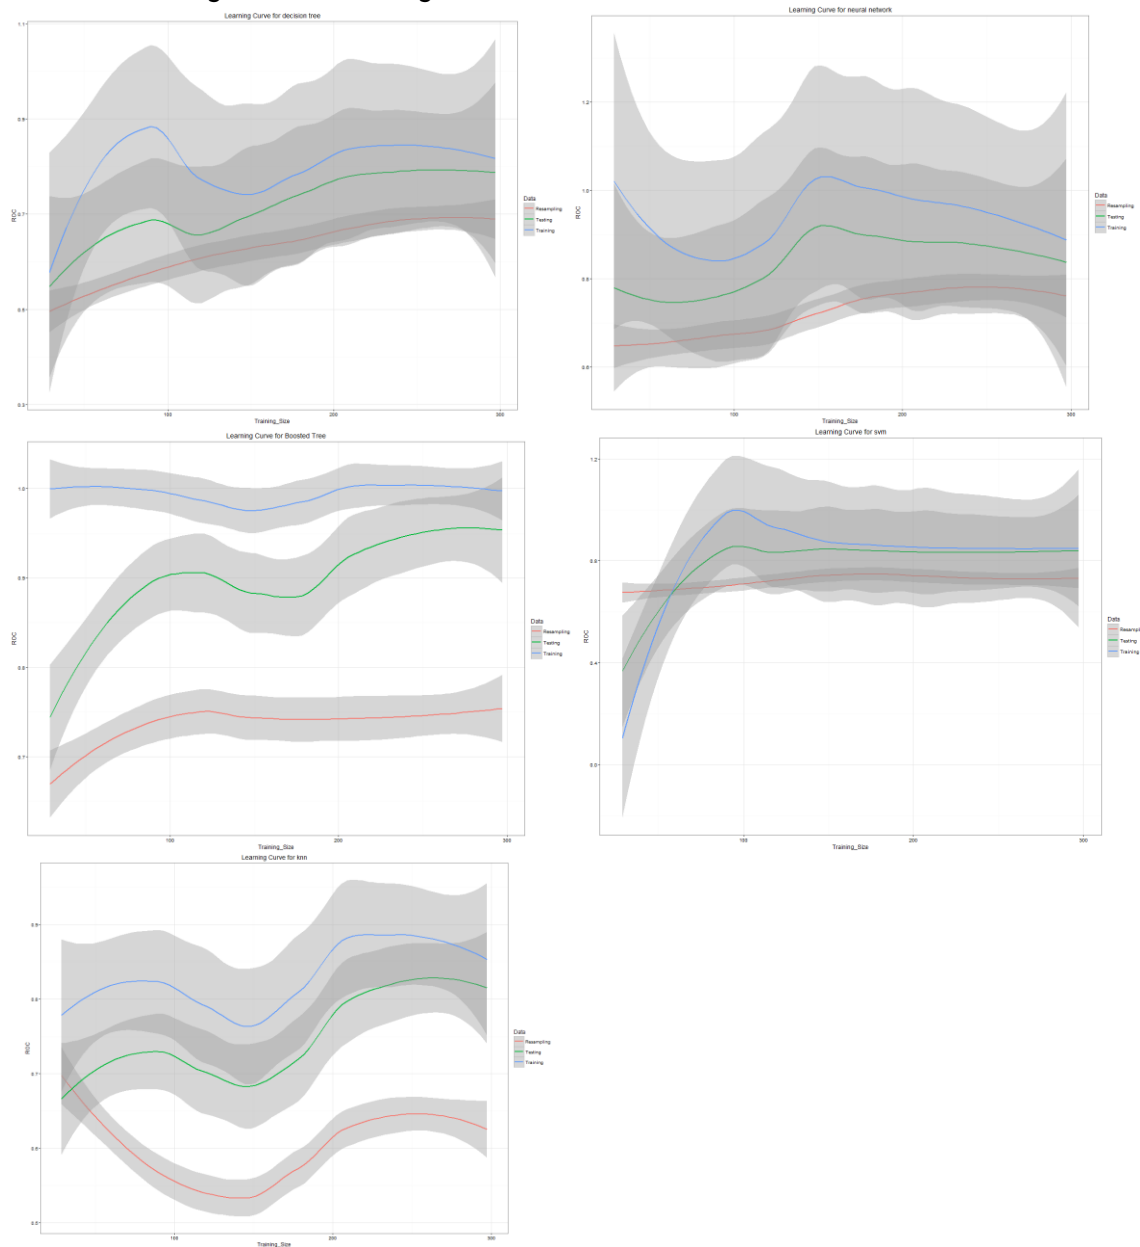


Fig.1 Learning Curves of 5 learning algorithms

Then we will randomly pick TrainData and TestData. We use histogram of TrainData and TestData to ensure they are similarly distributed. For example we could look at the age distribution in both training and testing.




Fig.2: Example distribution of train set




Fig.3: Example distribution of test set

3.3. Modeling learning algorithm

3.3.1. Decision Tree with some form of pruning.

Here we use package *tree* to do the decision tree model as it provide convenient cross validation as well pruning features.

The key code in r to build decision tree mode is,

```
treeMod <- tree(Consumption_additive_f ~ school + sex + age + address + famsize +
Pstatus + Medu + Fedu + Mjob + Fjob + traveltime + studytime + failures + higher
+ internet + romantic + famrel + freetime + goout + health + absences,
data=TrainData)
```

And then we could call “plot(treeMod); text(treeMod,pretty=0)” to plot the decision tree. As we can see from the below, we have 20 terminal nodes.




Fig.4 Decision Tree for Dataset 1




Fig.5: Pruning - misclass vs tree nodes

We could take a close look at the decision tree, it makes sense with our normal impression on how much a student drink to do with their behaviors or families. For example, students who go out often may drink more than those mostly stay at home. Or another example could be, male students tend to drink than females. Or the students who study more may not have that much time to drink alcohol as much as those who spend less time in study. This is to do with our domain knowledge, which makes sense as well. We can see we have 20 terminal nodes and we may have some overfitting situation here.

We could do some pruning by easily calling *cv.tree* function in package *tree*. Figure 5 lower X axis is the count of terminal nodes and upper X axis is the count of folds, i.e. # of pieces the data is split in the cross validation. This shows how this misclassification error average against the terminal node counts and number of folds. Hence the optimal number of terminal nodes should be one which gives lowest misclass while

smallest number of terminal nodes since usually the fewer nodes, the simpler decision tree model is. In this case, we can try best terminal nodes count = 5. Thus we could get the plot of the pruned decision tree as below.




Fig.6: Pruned Decision Tree

Because we do pruning, we are avoiding overfitting in the decision tree. Hence it is expected that we will lose some classification accuracy here using pruning for Training Dataset. But we should expect to improve accuracy in Test Dataset.

Table 1: Error rate for decision tree

	Without Pruning	With Pruning
Training set	0.1139241	0.1360759
Testing set	0.1898734	0.1772152

From the testing dataset, we can see the misclassification error reduces from 19% to 17.7%, which implies we did avoid overfitting and improve the generalization procedure for pruning.

3.3.2. Neural Networks

Here We use package *neuralnet* and need to install it at the first time by running `install.packages("neuralnet")`.

Step1: At the first step, we need to normalize the data before training a neural network. Here I used min-max scale method to scale the data to interval of [0,1].

Step2: We are going to build neural network model. Here we use logistic function as the activation function as we have classification at the end. Here we have 22 input at the first layer and target to have 1 output at the end. Usually we target to have 2/3 nodes of previous layers so we decide to have 10 and 5 layers as hidden layers. I would like to have activation function been applied to the output thus the `linear.output` set to FALSE.

```
nn <- neuralnet(Consumption_additive ~ school.f + sex.f + age + address.f +
  famsize.f + Pstatus.f + Medu + Fedu + Mjob.f + Fjob.f + traveltime + studytime +
  failures + higher.f + internet.f + romantic.f + famrel + freetime + goout +
  health + absences + Grades, data=TrainData, act.fct = "logistic", hidden =
  c(10,5), linear.output=FALSE)
```

And we could try different layers and nodes as well as the activation function to see the best neural network model.

Table 2: Best neural network model

Act.function	Layers	TrainErrorRate	TestErrorRate
logistic	c(10)	0.01898734177	0.2278481013
logistic	c(10,5)	0.01898734177	0.1392405063
logistic	c(15,10,5)	0.03164556962	0.1772151899
tanh	c(10,5)	0.01898734177	0.2278481013

And we could plot the neural network as below.

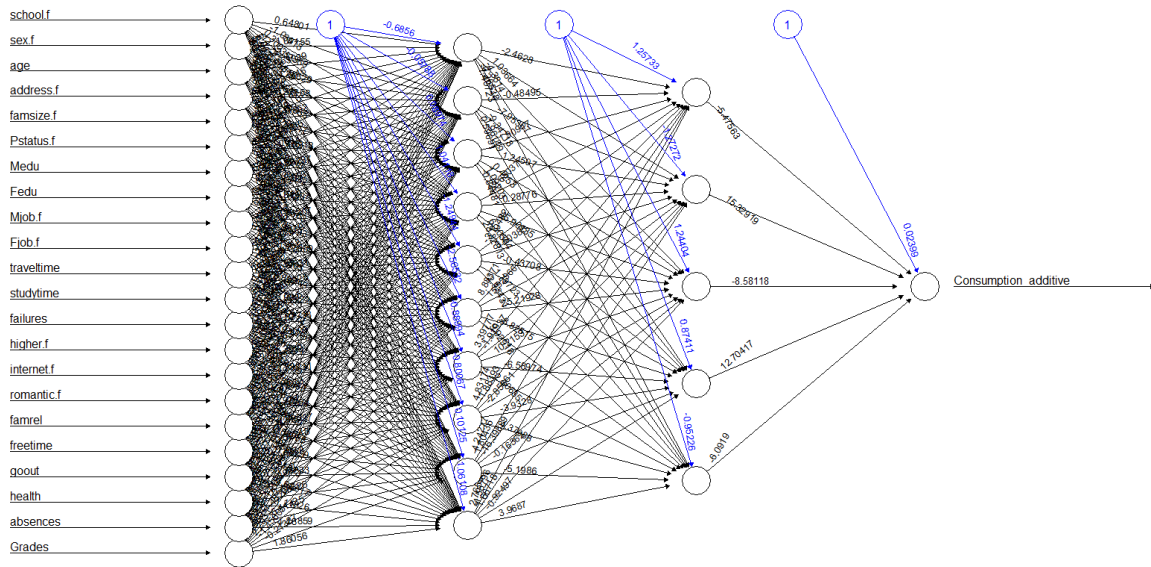


Fig.7: Neural Network plot for dataset #1

Step3: We check the accuracy of this neural network model.

Table 3: Neural Network error rate

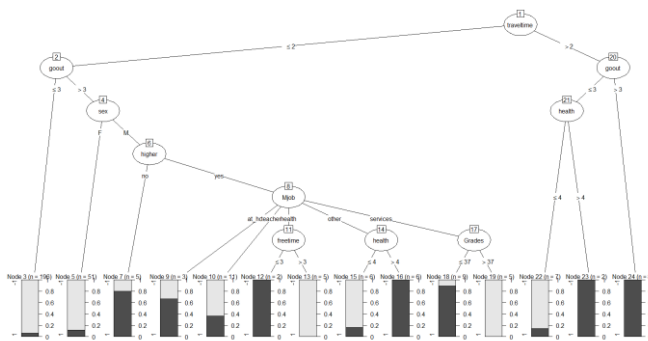
	error rate	error rate avg. from 10 times
Training set	0.01898734177	0.02974683544
Testing set	0.1392405063	0.2113924051

As we can see, the misclassification in training dataset is only 1.9% and the misclassification in test dataset is 13.9%. Comparing to Decision tree, which is 17~18% misclassification, here the neural network model we have built here is better.

Step4: Cross Validation by running 10 times of training and testing dataset split and then take the average misclassification error. With 10 times repeated run, we can get an idea how the neural model performs over other learning algorithm. We get average test dataset misclassification error about 21%, which is a bit worse than decision tree we have. Also the training set data error is only 0.03, which leads us to think neural network here may be overfitting already. But we have to admit that we have tried many combination of different layers and nodes in neural networks. However neural networks seems performs a bit slow in this case.

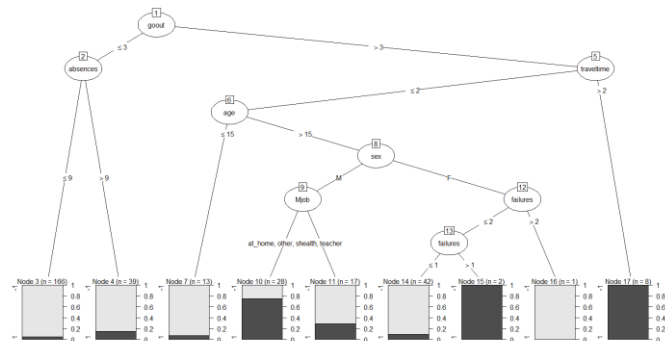
3.3.3. Boosting.

Here we use package C50 to do decision boosting as it provide a more convenient way to setting the boosting parameter.



Testing Error rate= 0.2278481

Fig. 8: Decision tree with some Pruning



Testing Error rate= 0.1898734

Fig.9: Boosting with Pruning

Firstly we do a decision tree again using C50 with some pruning. Then use Boost for the decision tree by assigning trials = 10. Also try to use feature selection (winnow=TRUE) to do more aggressive in pruning.

```
treeMod_boosting10 <- C5.0(TrainData[-23],
  TrainData$Consumption_additive.f, trials = 10, control =
  C5.0Control(noGlobalPruning = FALSE, earlyStopping = TRUE, CF=0.25, winnow=FALSE))
```

Hence the boost improves the accuracy by reducing error rate from 22.78% to 19%.

3.3.4. Support Vector Machines

We use package caret to do SVM. We use radial basic function kernel at the first and then use linear kernel. As always we will need to scale the data to regularize the values. From lecture, we know the larger C, the more fit in training set, which may leads to over fit the training set. Thus we will watch the error of test set to locate the best C.

Gaussian kernel

here we will try different C to find best model. We will do 10 folds and repeat 5 times for cross validation.

```
> ctrl <- trainControl(method="repeatedcv", # 10 fold cross validation
  repeats=5, # do 5 repetitions of cv
  summaryFunction=twoClassSummary,
  classProbs=TRUE) # Use AUC to pick the best model #
> svm.tune1 <- train(x=TrainX,y= TrainY,method = "svmRadial", # Radial kernel
+ tuneLength = 9, # 9 values of the cost function
+ preProc = c("center","scale"), # Center and scale data
+ metric="ROC", trControl=ctrl)
```

Hence for different C, we can pick the C which gives largest ROC. The final values used for the model were sigma = 0.02648622 and C = 1. Resampling results across tuning parameters:

C	ROC	Sens	Spec
0.25	0.8074256	0.9549231	0.3006667
0.50	0.8075333	0.9518769	0.3106667
1.00	0.8079744	0.9556923	0.2786667
2.00	0.7979231	0.9604615	0.2580000
4.00	0.7848410	0.9581231	0.2173333
8.00	0.7725179	0.9635077	0.1760000
16.00	0.7591333	0.9666462	0.1626667
32.00	0.7586718	0.9666462	0.1466667
64.00	0.7592513	0.9659077	0.1706667

Similarly we step further to fine tune to find best C and sigma for our model by giving C a range c(0.75, 0.9, 1, 1.1, 1.25) around 1 and sigma range c(.01, 0.0264, 0.3). Similarly we find the best C = 1.25 and sigma = 0.01. Then the train set and testing set error is as below.

Table 4: SVM error rate

	error rate of Radial kernel (Gaussian)	error rate of Linear kernel (no kernel)
Training set	0.1234177	0.1202532
Testing set	0.164557	0.1898734

Hence the misclassification error of training set is around 0.123 while misclassification error of test dataset is 0.164.

Linear Kernel

Linear kernel is similar process with Gaussian kernel except to explicitly define it by method = svmLinear in the key train function. We try different C as well by giving tuneGrid c(0.9, 1, 1.1, 1.25, 1.5, 2)

```
> grid2<- expand.grid(C = c(0.9, 1, 1.1, 1.25, 1.5, 2))
> svm.tune2 <- train(x=TrainX,y= TrainData$Consumption_additive,method =
"svmLinear",preProc = c("center","scale"),metric="ROC",tuneGrid =
grid2,trControl=ctrl)
```

Similarly ROC was used to select the optimal model using the largest value. The final value used for the model was C = 1.1. The training and testing error rate is in Table X. Hence the misclassification error of training set is around 0.12 while misclassification error of test dataset is 0.189.




Fig. 10: SVM kernel function ROC comparison

Therefore we can see the Gaussian kernel performs a little better accuracy than linear kernel for testing dataset. It seems that linear kernel here overfits a bit in trainset as linear kernel provides a bit better accuracy in training set than Gaussian kernel. As to the processing time, both kernel performs similarly quick as we don't have big size of data in this problem. And it did performs faster than neural network.

3.3.5. k-Nearest Neighbors.

Here we use package *class* to do knn learning. The key function to do knn algorithm in r is as below by using package *class*.

```
pred_test_knn = knn(TrainData[,-23], TestData[,-23], cl=TrainData[,23], k, l = 0,
prob = FALSE, use.all = TRUE)
pred_train_knn = knn.cv(TrainData, cl=TrainData[,23], k, l = 0, prob = FALSE,
use.all = TRUE)
```

Table 5: KNN error rate

k value	TrainDataErr	TestDataErr	k value	TrainDataErr	TestDataErr
1	0.136076	0.227848	7	0.164557	0.177215
2	0.167722	0.291139	8	0.167722	0.164557
3	0.129747	0.253165	9	0.177215	0.177215
4	0.148734	0.253165	10	0.180380	0.164557
5	0.161392	0.189873	11	0.174051	0.164557
6	0.170886	0.164557	12	0.183544	0.164557

Hence we can pick can k = 6 which can give us best Test data error.

4. Second dataset #2 - adult dataset

4.1. Data attributes

age	continuous.
workclass	Private, Self-emp-not-inc, Self-emp-inc, Federal, Local, State, Wo pay, Never-worked.
fnlwgt	continuous.
education	Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
education-num	continuous.
marital-status	Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
occupation	Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport etc.
relationship	Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race	White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
sex	Female, Male.
capital-gain	continuous.
capital-loss	continuous.
hours-per-week	continuous.
native-country	United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China etc.

4.2. Learning curve to help decide train data size

As we did previously, we take a look at the learning curve to see what the minimum training size for each algorithm is. As we could notice from the below that neural networks require least amount of training size while boosting and knn will require most training set size. For sanity to obtaining best model, we later randomly assign 80% of all data as training set and remaining 20% for test set.

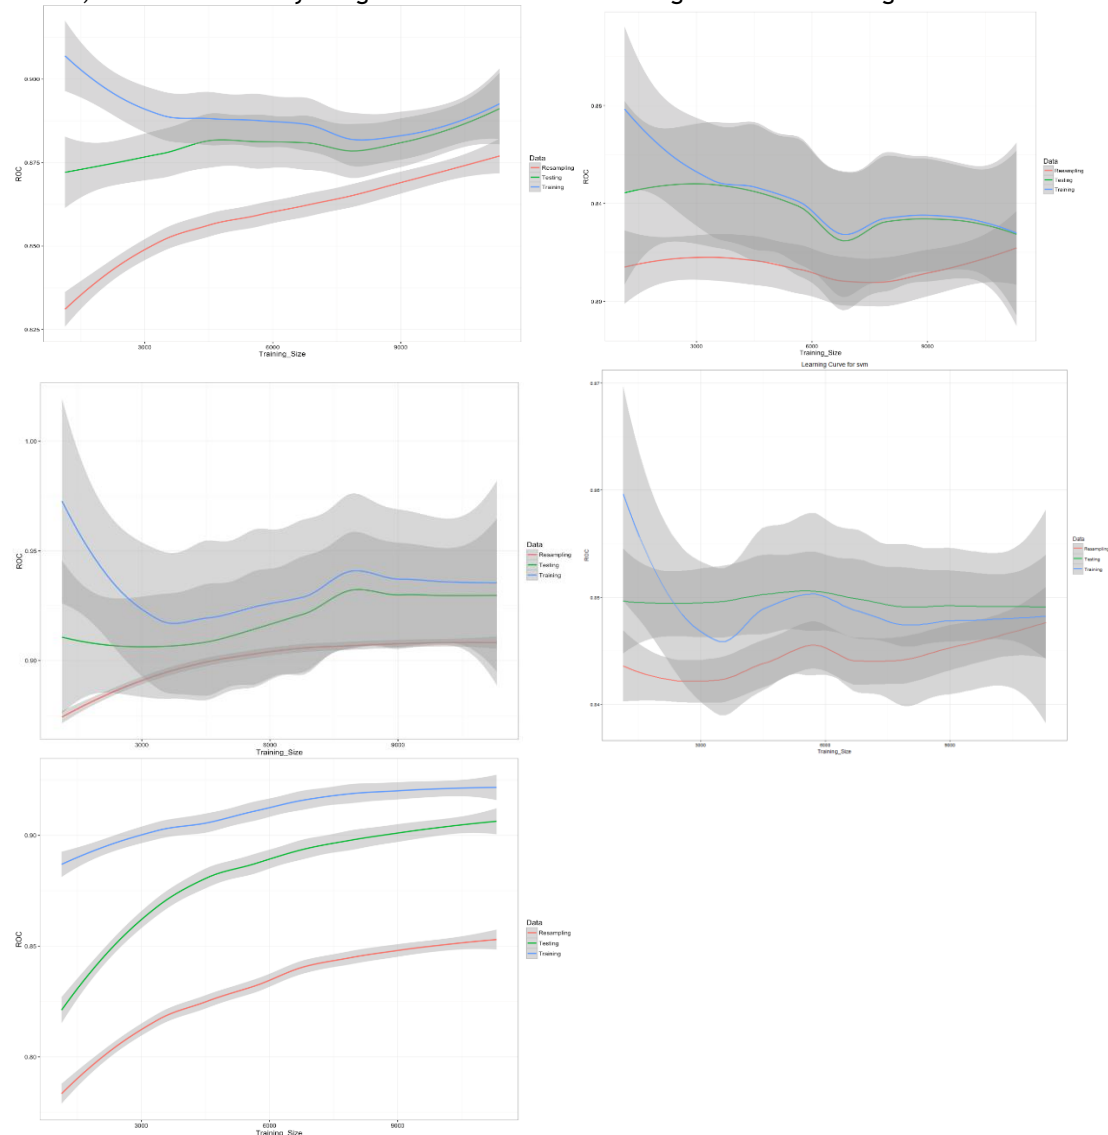


Fig. 11: Learning Curves for 5 algorithms (upper left right to bottom: decision tree, neural networks, boosting, svm, knn)

Similarly we will ensure randomly selected train and test set are similar distributed set. For example we could look at the age distribution of training and testing set.




Fig.12: Example distribution of train set




Fig.13: Example distribution of test set

4.3. Modeling learning algorithm

4.3.1. Decision tree

Similarly we use package tree to do decision tree algorithm.




Fig. 14: Decision Tree without pruning for Dataset #2




Fig.15: Pruning - misclass vs tree nodes

Table 6: Error rate for decision tree

	Without Pruning	With Pruning
Training set	0.157031	0.157031
Testing set	0.1665838	0.1665838

The cross validation and pruning Fig. tell us we should stay with minimum terminal nodes 5. Thus we will have the same training and testing error. The tree could be concise a little bit as the ending branch at the first one had same outputs. Thus it will look like below.

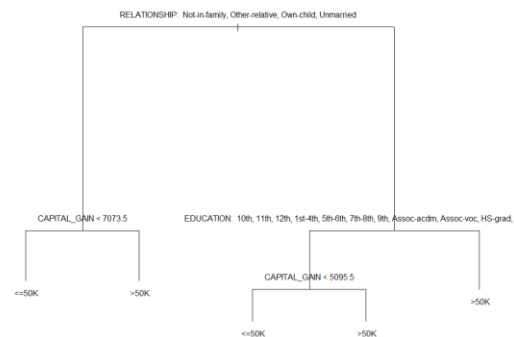


Fig. 16: Pruned Decision tree for dataset #2

4.3.2. Neural Network

As always, we will firstly normalize the data and I use max-min method to save it to [0,1] like previous dataset. We also use kernlab to se neuralnet function. The neural network model is very slow here due to large amount of data size. Thus I just do one try with one single hidden layer with 4 nodes and activation function tanh.

```
nn <- neuralnet(SALARY_CLASS_n ~ AGE + WORKCLASS_n + FNLWGT + EDUCATION_n +
  EDUCATION_NUM + MARITAL_STATUS_n + OCCUPATION_n + RELATIONSHIP_n + RACE_n + SEX_n
  + CAPITAL_GAIN + CAPITAL_LOSS + HOURS_PER_WEEK + NATIVE_COUNTRY_n,
  data=TrainData, act.fct = "tanh", hidden = 4, linear.output=T)
```

The cross validation was done by repeating the same process 10 times and the average error we got is around 68%. The error rate here is very big, I believe it is due to I have not tried to select a best hidden layers and the node numbers of each layer.

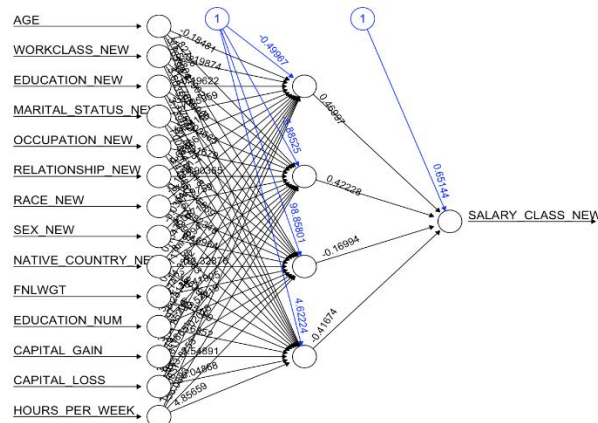


Fig. 17: Neural network plot for dataset#2

4.3.3. Boosting

Here similar to the first dataset, we use package C50 to do decision boosting as it provide a more convenient way to setting the boosting parameter.

Firstly we build the decision tree using C50 without pruning. The code to build decision tree model using C50 is as below

```
treeMod <- c5.0(TrainData[-15], TrainData$SALARY_CLASS_f, control =
  C5.0Control(noglobalPruning = TRUE, earlyStopping = FALSE, CF=0.25, winnow=FALSE))
```

actual default	predicted default		Row Total	actual default	predicted default		Row Total	actual default	predicted default		Row Total
	<=50K	>50K			<=50K	>50K			<=50K	>50K	
<=50K	4232	336	4568	<=50K	4237	331	4568	<=50K	4231	337	4568
	0.701	0.056			0.702	0.055			0.701	0.056	
>50K	490	975	1465	>50K	492	973	1465	>50K	483	982	1465
	0.081	0.162			0.082	0.161			0.080	0.163	
Column Total	4722	1311	6033	Column Total	4729	1304	6033	Column Total	4714	1319	6033

Test data error rate = 0.1369136416 Test data error rate = 0.1364163766 Test data error rate = 0.135919111

Fig.18: Decision Tree w/o pruning Fig.19: Decision Tree with pruning Fig.20: Decision tree with boosting

As to do more aggressively in pruning, we could play with winnow option. However the winnow option sometimes could help improve accuracy sometimes it does not. Because the winnow option is erroneously removing predictors that can improve the accuracy of the model. Within the cross-validation loop, the winnowing process thinks that it is improving the accuracy, but that is not holding up once other samples are used to evaluate performance. In this case (different than the first dataset), winnow option is not helping improving the accuracy.

4.3.4. SVM

Here we use Kernlab package and use function `ksvm` to build the SVM model. This function provides cross validation tuning as well as defining the C and sigma for the Gaussian kernel. Similar to what we have done before, we will try different parameters to find best model. Here we try different C for SVM functions. For computational sanity, we pick cross validation = 3, one can pick higher number such as 10 as we do before.

➤ Gaussian Kernel

For Gaussian Kernel, we use kernel = "rbfdot", then the key code to build SVM is,

```
svm_md1 <- ksvm(SALARY_CLASS_n~., data = TrainData, scaled = TRUE, kernel = "rbfdot",  
kpar=list(sigma=0.05), C=C1, cross=3)
```

The different C will gives the training set error and test error as below,

Table 7: SVM Gaussian error rate

C	TrainDataErr	TestDataErr	C	TrainDataErr	TestDataErr
0.25	0.158067	0.159456	8.00	0.139044	0.151500
0.50	0.154254	0.156473	16.00	0.134195	0.151997
1.00	0.150773	0.154649	32.00	0.129554	0.154318
2.00	0.147540	0.153158	64.00	0.123917	0.156141
4.00	0.143189	0.151832			

As we could see from different C, best C = 8 which gives lowest Test Data Error = 0.1515. Then we run with C = 8 again and we could see the cross validation error is 0.15545.

➤ Linear Kernel

For Linear Kernel, i.e. no kernel but just use x directly, we use kernel = "vanilladot", then the key code to build SVM is,

```
svm_md1 <- ksvm(SALARY_CLASS_n~., data = TrainData, scaled = TRUE,  
kernel="vanilladot", kpar=list(sigma=0.05), C=C1, cross=3)
```

Table 8: SVM Linear kernel error rate

C	TrainDataErr	TestDataErr	C	TrainDataErr	TestDataErr
0.25	0.18997886	0.19244157	4.00	0.18977164	0.19293884
0.50	0.19006175	0.19277308	8.00	0.18977164	0.19293884
1.00	0.18985453	0.19277308	16.00	0.18981309	0.19293884
2.00	0.18989598	0.19293884	32.00	0.18981309	0.19293884

Hence, best C = 1 which gives lowest Test Data Error = 0.19277. Then we run with C = 8 again and we could see the cross validation error is 0.190145.

As we can see the Gaussian kernel has better accuracy at test dataset as well as training set. Also a noticeable difference between kernel Gaussian and linear kernel is the processing time. The linear kernel took almost five times more clock time to process the grid of C parameter. The final C = 64 even take more than 30 mins and still it could not generate a result.

4.3.5. knn

Here similarly we use package `class` to do knn learning. The key function to do knn algorithm in r is as below by using package `class`.

```
pred_test_knn = knn(TrainData[, -15], TestData[, -15], cl=TrainData[, 15], k, l = 0,  
prob = FALSE, use.all = TRUE)  
pred_train_knn = knn.cv(TrainData[, -15], cl=TrainData[, 15], k, l = 0, prob =  
FALSE, use.all = TRUE)
```

Table 9: knn error rate

k value	TrainDataErr	TestDataErr	k value	TrainDataErr	TestDataErr
1	0.27825438	0.28692193	25	0.2087944	0.20968009
4	0.25355381	0.25178187	30	0.2090845	0.20984585

8	0.22379709	0.22708437	35	0.20954039	0.21100613
10	0.21642008	0.22028841	40	0.20999627	0.21183491
15	0.21016205	0.21034311	45	0.21140536	0.21249793
20	0.20858718	0.21084038	50	0.21198558	0.21382397

So the best $k=25$. The Test data error is 0.2097.

5. Conclusion

After running 5 learning algorithms, we could conduct a comparison between them to see the accuracy as well processing time needed.

Machine Learning in R	Decision Tree w/ Pruning	Neural Network	Boosting	SVM-Gaussian	SVM-Linear	kNN
Package used	tree	kernlab	C5.0	caret or ksvm	caret or ksvm	class
Dataset #1 rough est. Minimum Training size	300	400	500	150	150	600
Dataset #1 Training error rate	13.6%	3.0%	8.5%	12.3%	12.0%	17.1%
Dataset #1 Test error rate	17.7%	21.1%	19.0%	16.5%	19.0%	16.5%
Dataset #2 rough est. Minimum Training size	9000	3000	12000	10000	10000	16000
Dataset #2 Training error rate	15.7%	N/A	N/A	13.9%	19.0%	20.9%
Dataset #2 Test error rate	16.7%	68.0%	13.6%	15.2%	19.3%	21.0%
Processing speed	fast	median-slow	median	fast	fast	fast-median

As we can see from above, almost all the learning algorithms achieve similar accuracy in testing set, that is the error is around 15% ~ 20% for both datasets. The exception is the neural network for dataset # 2. As we have covered previously, I think this is due to the fact that I have not searched different hidden layers and nodes as computational time consuming. All the algorithms has done cross validation in the modelling.

For the decision tree of dataset #2, the pruning does not improved the accuracy in testing set because we are already at optimal terminal nodes counts.

Generally speaking, I think SVM Gaussian algorithm performs best from test set error rate and the processing speed point of view. And it seems to require least amount of training set.

6. References:

<https://cran.r-project.org/web/packages/C50/C50.pdf>
<http://www.cs.bc.edu/~alvarez/ML/statPruning.html>
<http://www.patricklamle.com/Tutorials/Decision%20tree%20R/Decision%20trees%20in%20R%20using%20C50.html>
<https://www.r-bloggers.com/the-5th-tribe-support-vector-machines-and-caret/>
<http://topepo.github.io/caret/miscellaneous-model-functions.html#yet-another-k-nearest-neighbor-function>
<https://cran.r-project.org/web/packages/class/class.pdf>